

12-Copy1

November 18, 2018

1 Assignment 12 - Neural Networks image recognition

Use both MLNN and the ConvNet to solve the following problem.

1. Add random noise (i.e. `np.random.normal`) to the images in training and testing. Make sure each image gets a different noise feature added to it. Inspect by printing out an image.
2. Compare the loss/accuracy (train, val) after N epochs for both MLNN and ConvNet with and without noise.
3. Vary the amount of noise (multiply `np.random.normal` by a factor) and keep track of the accuracy and loss (for training and validation) and plot these results.

2 Neural Networks - Image Recognition

```
In [1]: import keras
        from keras.datasets import mnist
        from keras.models import Sequential
        from keras.optimizers import RMSprop
        from keras.layers import Dense, Dropout, Flatten
        from keras.layers import Conv2D, MaxPooling2D
        from keras import backend
        from keras import utils as np_utils

        import numpy as np
        import matplotlib.pyplot as plt
        %matplotlib inline
```

Using TensorFlow backend.


2.1 Multi Layer Neural Network

Trains a simple deep NN on the MNIST dataset. Gets to 98.40% test accuracy after 20 epochs (there is *a lot* of margin for parameter tuning).

```
In [2]: # the data, shuffled and split between train and test sets
        (x_train, y_train), (x_test, y_test) = mnist.load_data()
```

```
In [3]: plt.imshow(x_train[0])
```

```
Out[3]: <matplotlib.image.AxesImage at 0x10ed86780>
```



```
In [4]: # Graph without noise
(x_train, y_train), (x_test, y_test) = mnist.load_data()

x_train = x_train.reshape(60000, 784)
x_test = x_test.reshape(10000, 784)
x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

batch_size = 128
num_classes = 10
epochs = 20

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

model = Sequential()
model.add(Dense(512, activation='relu', input_shape=(784,)))
```

```

model.add(Dropout(0.2))
model.add(Dense(512, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(10, activation='softmax'))

model.summary()

model.compile(loss='categorical_crossentropy',
              optimizer=RMSprop(),
              metrics=['accuracy'])

history = model.fit(x_train, y_train,
                   batch_size=batch_size,
                   epochs=epochs,
                   verbose=1,
                   validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

# Get training and test loss histories
training_loss = history.history['loss']
test_loss = history.history['val_loss']

# Create count of the number of epochs
epoch_count = range(1, len(training_loss) + 1)

# Visualize loss history
plt.plot(epoch_count, training_loss, 'r--')
plt.plot(epoch_count, test_loss, 'b-')
plt.legend(['Training Loss', 'Test Loss'])
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.show();

```

60000 train samples

10000 test samples


Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 512)	401920
dropout_1 (Dropout)	(None, 512)	0
dense_2 (Dense)	(None, 512)	262656
dropout_2 (Dropout)	(None, 512)	0

```

dense_3 (Dense)                (None, 10)                5130
=====
Total params: 669,706
Trainable params: 669,706
Non-trainable params: 0
-----
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 8s 126us/step - loss: 0.2479 - acc: 0.9236 - val_
Epoch 2/20
60000/60000 [=====] - 7s 118us/step - loss: 0.1023 - acc: 0.9695 - val_
Epoch 3/20
60000/60000 [=====] - 7s 120us/step - loss: 0.0753 - acc: 0.9774 - val_
Epoch 4/20
60000/60000 [=====] - 8s 125us/step - loss: 0.0596 - acc: 0.9826 - val_
Epoch 5/20
60000/60000 [=====] - 8s 126us/step - loss: 0.0510 - acc: 0.9842 - val_
Epoch 6/20
60000/60000 [=====] - 8s 126us/step - loss: 0.0437 - acc: 0.9865 - val_
Epoch 7/20
60000/60000 [=====] - 9s 142us/step - loss: 0.0387 - acc: 0.9883 - val_
Epoch 8/20
60000/60000 [=====] - 9s 157us/step - loss: 0.0334 - acc: 0.9902 - val_
Epoch 9/20
60000/60000 [=====] - 9s 144us/step - loss: 0.0316 - acc: 0.9909 - val_
Epoch 10/20
60000/60000 [=====] - 9s 147us/step - loss: 0.0310 - acc: 0.9910 - val_
Epoch 11/20
60000/60000 [=====] - 10s 165us/step - loss: 0.0268 - acc: 0.9924 - val_
Epoch 12/20
60000/60000 [=====] - 9s 151us/step - loss: 0.0239 - acc: 0.9930 - val_
Epoch 13/20
60000/60000 [=====] - 8s 136us/step - loss: 0.0249 - acc: 0.9929 - val_
Epoch 14/20
60000/60000 [=====] - 7s 124us/step - loss: 0.0217 - acc: 0.9939 - val_
Epoch 15/20
60000/60000 [=====] - 8s 125us/step - loss: 0.0231 - acc: 0.9937 - val_
Epoch 16/20
60000/60000 [=====] - 7s 122us/step - loss: 0.0199 - acc: 0.9941 - val_
Epoch 17/20
60000/60000 [=====] - 7s 122us/step - loss: 0.0186 - acc: 0.9949 - val_
Epoch 18/20
60000/60000 [=====] - 7s 123us/step - loss: 0.0191 - acc: 0.9945 - val_
Epoch 19/20
60000/60000 [=====] - 7s 123us/step - loss: 0.0176 - acc: 0.9953 - val_
Epoch 20/20
60000/60000 [=====] - 7s 122us/step - loss: 0.0179 - acc: 0.9953 - val_
Test loss: 0.10639518911132022

```

Test accuracy: 0.9837




```
In [ ]: # 98.37% accuracy
```

```
In [5]: # Noise sigma = 10
(x_train, y_train), (x_test, y_test) = mnist.load_data()

# x_train values after the noise
mu, sigma = 0, 10
noise = np.random.normal(mu, sigma, [28,28])
x_train = x_train - noise

plt.imshow(x_train[0])
```

```
Out[5]: <matplotlib.image.AxesImage at 0xb28b30d68>
```



```
In [6]: x_train = x_train.reshape(60000, 784)
        x_test = x_test.reshape(10000, 784)
        x_train = x_train.astype('float32')
        x_test = x_test.astype('float32')
        x_train /= 255
        x_test /= 255

        # convert class vectors to binary class matrices
        y_train = keras.utils.to_categorical(y_train, num_classes)
        y_test = keras.utils.to_categorical(y_test, num_classes)

        batch_size = 128
        num_classes = 10
        epochs = 20

        model = Sequential()
        model.add(Dense(512, activation='relu', input_shape=(784,)))
        model.add(Dropout(0.2))
        model.add(Dense(512, activation='relu'))
        model.add(Dropout(0.2))
        model.add(Dense(10, activation='softmax'))

        model.summary()

        model.compile(loss='categorical_crossentropy',
                      optimizer=RMSprop(),
```

```

metrics=['accuracy'])

history = model.fit(x_train, y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1,
                    validation_data=(x_test, y_test))

# Get training and test loss histories
training_loss = history.history['loss']
test_loss = history.history['val_loss']

# Create count of the number of epochs
epoch_count = range(1, len(training_loss) + 1)

# Visualize loss history
plt.plot(epoch_count, training_loss, 'r--')
plt.plot(epoch_count, test_loss, 'b-')
plt.legend(['Training Loss', 'Test Loss'])
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.show();

```

Layer (type)	Output Shape	Param #
dense_4 (Dense)	(None, 512)	401920
dropout_3 (Dropout)	(None, 512)	0
dense_5 (Dense)	(None, 512)	262656
dropout_4 (Dropout)	(None, 512)	0
dense_6 (Dense)	(None, 10)	5130

```

Total params: 669,706
Trainable params: 669,706
Non-trainable params: 0

```

```

Train on 60000 samples, validate on 10000 samples

```

```
Epoch 1/20
```

```
60000/60000 [=====] - 8s 131us/step - loss: 0.2603 - acc: 0.9196 - val_
```

```
Epoch 2/20
```

```
60000/60000 [=====] - 7s 119us/step - loss: 0.1087 - acc: 0.9670 - val_
```

```
Epoch 3/20
```


```
60000/60000 [=====] - 7s 122us/step - loss: 0.0806 - acc: 0.9760 - val_
```

```
Epoch 4/20
```

```

60000/60000 [=====] - 7s 123us/step - loss: 0.0648 - acc: 0.9815 - val_
Epoch 5/20
60000/60000 [=====] - 8s 127us/step - loss: 0.0551 - acc: 0.9834 - val_
Epoch 6/20
60000/60000 [=====] - 8s 127us/step - loss: 0.0472 - acc: 0.9857 - val_
Epoch 7/20
60000/60000 [=====] - 7s 123us/step - loss: 0.0427 - acc: 0.9874 - val_
Epoch 8/20
60000/60000 [=====] - 8s 128us/step - loss: 0.0383 - acc: 0.9888 - val_
Epoch 9/20
60000/60000 [=====] - 7s 121us/step - loss: 0.0355 - acc: 0.9898 - val_
Epoch 10/20
60000/60000 [=====] - 7s 121us/step - loss: 0.0352 - acc: 0.9904 - val_
Epoch 11/20
60000/60000 [=====] - 7s 121us/step - loss: 0.0336 - acc: 0.9908 - val_
Epoch 12/20
60000/60000 [=====] - 7s 121us/step - loss: 0.0308 - acc: 0.9919 - val_
Epoch 13/20
60000/60000 [=====] - 7s 122us/step - loss: 0.0305 - acc: 0.9920 - val_
Epoch 14/20
60000/60000 [=====] - 7s 123us/step - loss: 0.0291 - acc: 0.9923 - val_
Epoch 15/20
60000/60000 [=====] - 7s 123us/step - loss: 0.0268 - acc: 0.9931 - val_
Epoch 16/20
60000/60000 [=====] - 8s 126us/step - loss: 0.0272 - acc: 0.9930 - val_
Epoch 17/20
60000/60000 [=====] - 8s 131us/step - loss: 0.0253 - acc: 0.9935 - val_
Epoch 18/20
60000/60000 [=====] - 8s 126us/step - loss: 0.0269 - acc: 0.9934 - val_
Epoch 19/20
60000/60000 [=====] - 7s 125us/step - loss: 0.0284 - acc: 0.9934 - val_
Epoch 20/20
60000/60000 [=====] - 8s 130us/step - loss: 0.0264 - acc: 0.9938 - val_


```

```
In [ ]: # 96.91 percent accuracy
```

```
In [8]: (x_train, y_train), (x_test, y_test) = mnist.load_data()  
mu, sigma = 0, 50  
noise = np.random.normal(mu, sigma, [28,28])  
x_train = x_train - noise  
x_train[0]  
plt.imshow(x_train[0])
```

```
Out[8]: <matplotlib.image.AxesImage at 0xb29240e10>
```



```
In [9]: x_train = x_train.reshape(60000, 784)
x_test = x_test.reshape(10000, 784)
x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

batch_size = 128
num_classes = 10
epochs = 20

model = Sequential()
model.add(Dense(512, activation='relu', input_shape=(784,)))
model.add(Dropout(0.2))
model.add(Dense(512, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(10, activation='softmax'))

model.summary()

model.compile(loss='categorical_crossentropy',
              optimizer=RMSprop(),
```

```

metrics=['accuracy'])

history = model.fit(x_train, y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1,
                    validation_data=(x_test, y_test))

# Get training and test loss histories
training_loss = history.history['loss']
test_loss = history.history['val_loss']

# Create count of the number of epochs
epoch_count = range(1, len(training_loss) + 1)

# Visualize loss history
plt.plot(epoch_count, training_loss, 'r--')
plt.plot(epoch_count, test_loss, 'b-')
plt.legend(['Training Loss', 'Test Loss'])
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.show();

```

Layer (type)	Output Shape	Param #
dense_7 (Dense)	(None, 512)	401920
dropout_5 (Dropout)	(None, 512)	0
dense_8 (Dense)	(None, 512)	262656
dropout_6 (Dropout)	(None, 512)	0
dense_9 (Dense)	(None, 10)	5130

```

Total params: 669,706
Trainable params: 669,706
Non-trainable params: 0

```

```

Train on 60000 samples, validate on 10000 samples

```

```
Epoch 1/20
```

```
60000/60000 [=====] - 9s 152us/step - loss: 0.3090 - acc: 0.9022 - val_
```

```
Epoch 2/20
```

```
60000/60000 [=====] - 8s 137us/step - loss: 0.1346 - acc: 0.9590 - val_
```

```
Epoch 3/20
```


```
60000/60000 [=====] - 7s 120us/step - loss: 0.1017 - acc: 0.9685 - val_
```

```
Epoch 4/20
```

```

60000/60000 [=====] - 7s 120us/step - loss: 0.0871 - acc: 0.9739 - val_
Epoch 5/20
60000/60000 [=====] - 7s 125us/step - loss: 0.0784 - acc: 0.9768 - val_
Epoch 6/20
60000/60000 [=====] - 7s 125us/step - loss: 0.0722 - acc: 0.9792 - val_
Epoch 7/20
60000/60000 [=====] - 8s 131us/step - loss: 0.0670 - acc: 0.9807 - val_
Epoch 8/20
60000/60000 [=====] - 8s 128us/step - loss: 0.0617 - acc: 0.9824 - val_
Epoch 9/20
60000/60000 [=====] - 8s 130us/step - loss: 0.0610 - acc: 0.9833 - val_
Epoch 10/20
60000/60000 [=====] - 8s 133us/step - loss: 0.0592 - acc: 0.9842 - val_
Epoch 11/20
60000/60000 [=====] - 8s 135us/step - loss: 0.0593 - acc: 0.9848 - val_
Epoch 12/20
60000/60000 [=====] - 8s 136us/step - loss: 0.0549 - acc: 0.9852 - val_
Epoch 13/20
60000/60000 [=====] - 8s 138us/step - loss: 0.0558 - acc: 0.9862 - val_
Epoch 14/20
60000/60000 [=====] - 8s 133us/step - loss: 0.0471 - acc: 0.9877 - val_
Epoch 15/20
60000/60000 [=====] - 8s 131us/step - loss: 0.0528 - acc: 0.9873 - val_
Epoch 16/20
60000/60000 [=====] - 8s 131us/step - loss: 0.0509 - acc: 0.9878 - val_
Epoch 17/20
60000/60000 [=====] - 8s 131us/step - loss: 0.0488 - acc: 0.9880 - val_
Epoch 18/20
60000/60000 [=====] - 8s 136us/step - loss: 0.0499 - acc: 0.9879 - val_
Epoch 19/20
60000/60000 [=====] - 8s 128us/step - loss: 0.0496 - acc: 0.9882 - val_
Epoch 20/20
60000/60000 [=====] - 8s 126us/step - loss: 0.0495 - acc: 0.9883 - val_


```



In []: # 95.64 percent accuracy

```
In [12]: (x_train, y_train), (x_test, y_test) = mnist.load_data()
mu, sigma = 0, 99
noise = np.random.normal(mu, sigma, [28,28])
x_train = x_train - noise
x_train[0]
plt.imshow(x_train[0])
```

Out[12]: <matplotlib.image.AxesImage at 0xb2a907518>



```
In [13]: x_train = x_train.reshape(60000, 784)
x_test = x_test.reshape(10000, 784)
x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

batch_size = 128
num_classes = 10
epochs = 20

model = Sequential()
model.add(Dense(512, activation='relu', input_shape=(784,)))
model.add(Dropout(0.2))
model.add(Dense(512, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(10, activation='softmax'))

model.summary()

model.compile(loss='categorical_crossentropy',
              optimizer=RMSprop(),
```

```

metrics=['accuracy'])

history = model.fit(x_train, y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1,
                    validation_data=(x_test, y_test))

# Get training and test loss histories
training_loss = history.history['loss']
test_loss = history.history['val_loss']

# Create count of the number of epochs
epoch_count = range(1, len(training_loss) + 1)

# Visualize loss history
plt.plot(epoch_count, training_loss, 'r--')
plt.plot(epoch_count, test_loss, 'b-')
plt.legend(['Training Loss', 'Test Loss'])
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.show();

```

Layer (type)	Output Shape	Param #
dense_10 (Dense)	(None, 512)	401920
dropout_7 (Dropout)	(None, 512)	0
dense_11 (Dense)	(None, 512)	262656
dropout_8 (Dropout)	(None, 512)	0
dense_12 (Dense)	(None, 10)	5130

```

Total params: 669,706
Trainable params: 669,706
Non-trainable params: 0

```

```

Train on 60000 samples, validate on 10000 samples

```

```
Epoch 1/20
```

```
60000/60000 [=====] - 8s 134us/step - loss: 0.4268 - acc: 0.8630 - val_
```

```
Epoch 2/20
```

```
60000/60000 [=====] - 7s 116us/step - loss: 0.1871 - acc: 0.9421 - val_
```

```
Epoch 3/20
```


```
60000/60000 [=====] - 7s 119us/step - loss: 0.1456 - acc: 0.9556 - val_
```

```
Epoch 4/20
```

```

60000/60000 [=====] - 8s 126us/step - loss: 0.1228 - acc: 0.9621 - val_
Epoch 5/20
60000/60000 [=====] - 8s 126us/step - loss: 0.1105 - acc: 0.9673 - val_
Epoch 6/20
60000/60000 [=====] - 8s 127us/step - loss: 0.1029 - acc: 0.9692 - val_
Epoch 7/20
60000/60000 [=====] - 8s 128us/step - loss: 0.0985 - acc: 0.9711 - val_
Epoch 8/20
60000/60000 [=====] - 8s 129us/step - loss: 0.0951 - acc: 0.9729 - val_
Epoch 9/20
60000/60000 [=====] - 8s 129us/step - loss: 0.0914 - acc: 0.9740 - val_
Epoch 10/20
60000/60000 [=====] - 8s 129us/step - loss: 0.0873 - acc: 0.9756 - val_
Epoch 11/20
60000/60000 [=====] - 8s 128us/step - loss: 0.0864 - acc: 0.9756 - val_
Epoch 12/20
60000/60000 [=====] - 8s 132us/step - loss: 0.0865 - acc: 0.9764 - val_
Epoch 13/20
60000/60000 [=====] - 8s 133us/step - loss: 0.0878 - acc: 0.9761 - val_
Epoch 14/20
60000/60000 [=====] - 8s 133us/step - loss: 0.0839 - acc: 0.9776 - val_
Epoch 15/20
60000/60000 [=====] - 8s 135us/step - loss: 0.0849 - acc: 0.9777 - val_
Epoch 16/20
60000/60000 [=====] - 8s 133us/step - loss: 0.0819 - acc: 0.9783 - val_
Epoch 17/20
60000/60000 [=====] - 8s 135us/step - loss: 0.0824 - acc: 0.9788 - val_
Epoch 18/20
60000/60000 [=====] - 8s 129us/step - loss: 0.0816 - acc: 0.9792 - val_
Epoch 19/20
60000/60000 [=====] - 8s 134us/step - loss: 0.0858 - acc: 0.9794 - val_
Epoch 20/20
60000/60000 [=====] - 8s 128us/step - loss: 0.0844 - acc: 0.9795 - val_

```

2.2 Conv Net

Trains a simple convnet on the MNIST dataset. Gets to 99.25% test accuracy after 12 epochs (there is still a lot of margin for parameter tuning).

```
In [14]: # input image dimensions
img_rows, img_cols = 28, 28

# the data, shuffled and split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()

if backend.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
    x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
    input_shape = (1, img_rows, img_cols)
else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
    x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
    input_shape = (img_rows, img_cols, 1)

x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print('x_train shape:', x_train.shape)
```

```

print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

```

```

x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples

```

```

In [15]: # convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

```

```

In [18]: batch_size = 128
num_classes = 10
epochs = 2

```

```

model = Sequential()
model.add(Conv2D(32, kernel_size=(3, 3),
                 activation='relu',
                 input_shape=input_shape))
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(num_classes, activation='softmax'))

model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adadelta(),
              metrics=['accuracy'])

history = model.fit(x_train, y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1,
                    validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

Train on 60000 samples, validate on 10000 samples

Epoch 1/2

60000/60000 [=====] - 151s 3ms/step - loss: 0.2123 - acc: 0.9348 - val_

Epoch 2/2

60000/60000 [=====] - 150s 2ms/step - loss: 0.0602 - acc: 0.9821 - val_

Test loss: 0.03919108632341958

Test accuracy: 0.9869


```

In [20]: # Get training and test loss histories
training_loss = history.history['loss']
test_loss = history.history['val_loss']

# Create count of the number of epochs
epoch_count = range(1, len(training_loss) + 1)

# Visualize loss history
plt.plot(epoch_count, training_loss, 'r--')
plt.plot(epoch_count, test_loss, 'b-')
plt.legend(['Training Loss', 'Test Loss'])
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.show();

```



```

In [22]: # input image dimensions
img_rows, img_cols = 28, 28


# the data, shuffled and split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()

mu, sigma = 0, 35
noise = np.random.normal(mu, sigma, [28,28])
x_train = x_train - noise

```

```
x_train[0]
plt.imshow(x_train[0])
```

Out[22]: <matplotlib.image.AxesImage at 0xb2e3fa2b0>



```
In [23]: if backend.image_data_format() == 'channels_first':
          x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
          x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
          input_shape = (1, img_rows, img_cols)
        else:
          x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
          x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
          input_shape = (img_rows, img_cols, 1)

          x_train = x_train.astype('float32')
          x_test = x_test.astype('float32')
          x_train /= 255
          x_test /= 255

          # convert class vectors to binary class matrices
          y_train = keras.utils.to_categorical(y_train, num_classes)
          y_test = keras.utils.to_categorical(y_test, num_classes)

          batch_size = 128
          num_classes = 10
          epochs = 2
```

```

model = Sequential()
model.add(Conv2D(32, kernel_size=(3, 3),
                 activation='relu',
                 input_shape=input_shape))
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(num_classes, activation='softmax'))

model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adadelta(),
              metrics=['accuracy'])

history = model.fit(x_train, y_train,
                   batch_size=batch_size,
                   epochs=epochs,
                   verbose=1,
                   validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

Train on 60000 samples, validate on 10000 samples

Epoch 1/2

60000/60000 [=====] - 161s 3ms/step - loss: 0.2315 - acc: 0.9286 - val_

Epoch 2/2

60000/60000 [=====] - 160s 3ms/step - loss: 0.0648 - acc: 0.9800 - val_

Test loss: 0.04118220631625736

Test accuracy: 0.9857

In [24]: *# Get training and test loss histories*

```
training_loss = history.history['loss']
```

```
test_loss = history.history['val_loss']
```

Create count of the number of epochs

```
epoch_count = range(1, len(training_loss) + 1)
```

Visualize loss history

```
plt.plot(epoch_count, training_loss, 'r--')
```


```
plt.plot(epoch_count, test_loss, 'b-')
```

```
plt.legend(['Training Loss', 'Test Loss'])
```

```
plt.xlabel('Epoch')
```

```
plt.ylabel('Loss')
```

```
plt.show();
```




```
In [25]: # input image dimensions
img_rows, img_cols = 28, 28

# the data, shuffled and split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()

mu, sigma = 0, 85
noise = np.random.normal(mu, sigma, [28,28])
x_train = x_train - noise
x_train[0]
plt.imshow(x_train[0])
```

```
Out[25]: <matplotlib.image.AxesImage at 0xb356d7eb8>
```



```
In [26]: if backend.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
    x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
    input_shape = (1, img_rows, img_cols)
else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
    x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
    input_shape = (img_rows, img_cols, 1)

x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

batch_size = 128
num_classes = 10
epochs = 8

model = Sequential()
model.add(Conv2D(32, kernel_size=(3, 3),
                 activation='relu',
                 input_shape=input_shape))
```

```

model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adadelta(),
              metrics=['accuracy'])

history = model.fit(x_train, y_train,
                   batch_size=batch_size,
                   epochs=epochs,
                   verbose=1,
                   validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

Train on 60000 samples, validate on 10000 samples

Epoch 1/8

60000/60000 [=====] - 169s 3ms/step - loss: 0.3801 - acc: 0.8812 - val_

Epoch 2/8

60000/60000 [=====] - 160s 3ms/step - loss: 0.1280 - acc: 0.9623 - val_

Epoch 3/8

60000/60000 [=====] - 163s 3ms/step - loss: 0.0966 - acc: 0.9716 - val_

Epoch 4/8

60000/60000 [=====] - 154s 3ms/step - loss: 0.0792 - acc: 0.9766 - val_

Epoch 5/8

60000/60000 [=====] - 151s 3ms/step - loss: 0.0707 - acc: 0.9789 - val_

Epoch 6/8

60000/60000 [=====] - 151s 3ms/step - loss: 0.0620 - acc: 0.9825 - val_

Epoch 7/8

60000/60000 [=====] - 153s 3ms/step - loss: 0.0560 - acc: 0.9828 - val_

Epoch 8/8

60000/60000 [=====] - 160s 3ms/step - loss: 0.0534 - acc: 0.9839 - val_

Test loss: 0.04573370984120675

Test accuracy: 0.9885

In [27]: # Get training and test loss histories

```
training_loss = history.history['loss']
```

```
test_loss = history.history['val_loss']
```

```
# Create count of the number of epochs
```

```
epoch_count = range(1, len(training_loss) + 1)
```



```
# Visualize loss history
plt.plot(epoch_count, training_loss, 'r--')
plt.plot(epoch_count, test_loss, 'b-')
plt.legend(['Training Loss', 'Test Loss'])
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.show();
```

