

Week 1 hello world

September 9, 2018

```
In [1]: import numpy as np
import pandas as pd
%matplotlib inline
import matplotlib.pyplot as plt
import os
```

```
In [2]: data = pd.DataFrame({'weight':[150, 100, 110, 122, 80, 90, 100, 180, 200, 100, 200, 120,
                                     'age'  :[ 17,   9,  13,  12,  6,  7,  12,  16,  17,  10,  20,  15,
data
```

```
Out[2]:
```

	weight	age
0	150	17
1	100	9
2	110	13
3	122	12
4	80	6
5	90	7
6	100	12
7	180	16
8	200	17
9	100	10
10	200	20
11	120	15
12	80	8
13	10	2
14	200	20


```
In [3]: data.describe()
```

```
Out[3]:
```

	weight	age
count	15.000000	15.000000
mean	122.800000	12.266667
std	54.173267	5.297798
min	10.000000	2.000000
25%	95.000000	8.500000
50%	110.000000	12.000000
75%	165.000000	16.500000
max	200.000000	20.000000

```
In [4]: data.plot(kind='scatter', x='weight', y='age', figsize=(5,5))
```

```
Out[4]: <matplotlib.axes._subplots.AxesSubplot at 0x115231048>
```



```
In [5]: def computeCost(X, y, theta):  
        inner = np.power(((X * theta.T) - y), 2)  
        return np.sum(inner) / (2 * len(X))
```

```
In [6]: data.insert(0, 'theOnes', 1)
```

```
In [7]: cols = data.shape[1]  
        X = data.iloc[:,0:cols-1]  
        y = data.iloc[:,cols-1:cols]
```

```
In [8]: print(X.head())  
        print(y.head())
```

```
   theOnes  weight  
0         1     150  
1         1     100  
2         1     110
```

```

3      1      122
4      1      80
    age
0      17
1       9
2      13
3      12
4       6

```

```

In [9]: X = np.matrix(X.values)
        y = np.matrix(y.values)
        theta = np.matrix(np.array([0,0]))
        theta

```

```

Out[9]: matrix([[0, 0]])

```

```

In [10]: X.shape, theta.shape, y.shape

```

```

Out[10]: ((15, 2), (1, 2), (15, 1))

```

```

In [11]: computeCost(X, y, theta)

```

```

Out[11]: 88.33333333333333

```

```

In [12]: def gradientDescent(X, y, theta, alpha, iters):
        temp = np.matrix(np.zeros(theta.shape))
        parameters = int(theta.ravel().shape[1])
        jhuAAP = np.zeros(iters)

        for i in range(iters):
            error = (X * theta.T) - y

            for j in range(parameters):
                term = np.multiply(error, X[:,j])
                temp[0,j] = theta[0,j] - ((alpha / len(X)) * np.sum(term))

            theta = temp
            jhuAAP[i] = computeCost(X, y, theta)

        return theta, jhuAAP

```

```

In [35]: alpha = 0.01
        iters = 67
        # Jim to Professor: Adding more iterations led to an error. Stackoverflow tells me this

```

```

In [36]: g, jhuAAP = gradientDescent(X, y, theta, alpha, iters)
        g

```

```

Out[36]: matrix([[3.01675102e+147, 4.37742847e+149]])

```

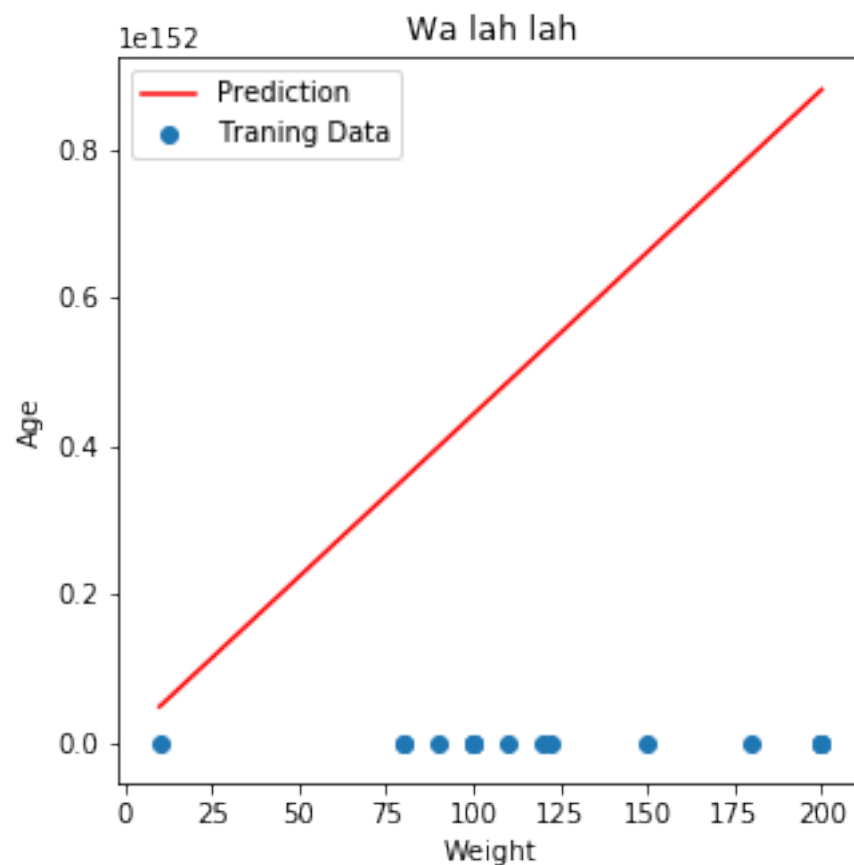
```
In [37]: computeCost(X, y, g)
```

```
Out[37]: 1.7073834821157e+303
```

```
In [47]: x = np.linspace(data.weight.min(), data.weight.max(), 80)
         f = g[0, 1] + (g[0, 1] * x)
```

```
fig, ax = plt.subplots(figsize=(5,5))
ax.plot(x, f, 'r', label='Prediction')
ax.scatter(data.weight, data.age, label='Traning Data')
ax.legend(loc=2)
ax.set_xlabel('Weight')
ax.set_ylabel('Age')
ax.set_title('Wa lah lah')
```

```
Out[47]: Text(0.5,1,'Wa lah lah')
```




```
In [39]: # Well, that's a fail.
```

```
In [48]: fig, ax = plt.subplots(figsize=(5,5))
         ax.plot(np.arange(iters), jhuAAP, 'r')
```

```
ax.set_xlabel('Iterations')
ax.set_ylabel('Age')
ax.set_title('Error vs. Training Epoch')
```

Out[48]: Text(0.5,1,'Error vs. Training Epoch')



In []: # Okay, I don't know what I'm doing.